

ANDREW ZHAO | Curriculum Vitae

Department of Automation, Tsinghua University – Citizenship: Canada

📞 (+86) 18513245436 • ✉ andrewzhao112@gmail.com

🌐 Personal Page 🎓 Google Scholar 🤖 Github 💬 LinkedIn 🐦 X/Twitter

💡 Research Statement

My research focuses on **reinforcement learning for LLMs**, **reasoning**, **LLM-based agents**, and the **safety** challenges associated with generative AI, with an emphasis on diversity, robustness, generalization, and self-evolving systems requiring minimal human intervention. Recently, I have become interested in the reasoning capabilities of large language models. I aim to advance the development of innovative, reliable, and autonomous AI systems capable of adapting to diverse and dynamic environments.

🎓 Education

Tsinghua University

Beijing

PhD, *Pattern Recognition and Machine Learning*, China

2021–present

Reinforcement Learning, Agents, Natural Language Processing, Deep Learning

Advisor :: Professor Gao Huang, Department of Automation

University of Southern California

Los Angeles

Master of Science, USA

2019–2020

University of British Columbia

Vancouver

Bachelor of Applied Science, Canada

2012–2017

📄 Selected Publications

Andrew Zhao, Yiran Wu, Yang Yue, Tong Wu, Quentin Xu, Yang Yue, Matthieu Lin, Shenzhi Wang, Qingyun Wu, Zilong Zheng, and Gao Huang. Absolute Zero: Reinforced Self-play Reasoning with Zero Data  * 1.8k. *Advances in Neural Information Processing Systems (NeurIPS)*  SPOTLIGHT, 2025.

Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu Lin, Y. Liu, and Gao Huang. ExpeL: LLM Agents Are Experiential Learners  * 188. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*  ORAL, 2024.

Yang Yue*, Zhiqi Chen*, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang. Does Reinforcement Learning Really Incentivize Reasoning Capacity in LLMs Beyond the Base Model? *Advances in Neural Information Processing Systems (NeurIPS)*  BEST PAPER RUNNER-UP, 2025.

Andrew Zhao, Quentin Xu, Matthieu Lin, Shenzhi Wang, Yongjin Liu, Zilong Zheng, and Gao Huang. DiveR-CT: Diversity-enhanced Red Teaming Large Language Model Assistants with Relaxing Constraints. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*  ORAL, 2025.

Haojun Jiang*, Andrew Zhao*, Qian Yang*, Xiangjie Yan, Teng Wang, Yulin Wang, Ning Jia, Jiangshan Wang, Guokun Wu, Yang Yue, Shaqi Luo, Huanqian Wang, Ling Ren, Siming Chen, Pan Liu, Guocai Yao, Wenming Yang, Shiji Song, Xiang Li, Kunlun He, and Gao Huang. Towards Expert-level Autonomous Carotid Ultrasonography with Large-scale Learning-based Robotic System. *Nature Communications*, 2025.

Andrew Zhao, Reshma Ghosh, Vitor Carvalho, Emily Lawton, Keegan Hines, Gao Huang, and Jack W. Stokes. Are My Optimized Prompts Compromised? Exploring Vulnerabilities of LLM-based Optimizers. *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2026.

Andrew Zhao*, Matthieu Lin*, Yangguang Li, Yong-Jin Liu, and Gao Huang. A Mixture of Surprises for Unsupervised Reinforcement Learning. *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Andrew Zhao*, Erle Zhu*, Rui Lu, Matthieu Lin, Yong-Jin Liu, and Gao Huang. Self-Referencing Agents for Unsupervised Reinforcement Learning. *Neural Networks*, page 107448, 2025.

Research Experience

Microsoft Research, Redmond

Research Intern @ AI Interaction and Learning Group

June, 2025 – August, 2025

Worked on AI Safety of Self-improving LLM Systems.

Advisor :: Dr. Jack Stokes, Principal RSDE @ MSR Redmond

Microsoft Research, Redmond

Research Intern @ Augmented Learning and Reasoning Group

June, 2024 – August, 2024

Worked on LLM-based agent applications for cybersecurity investigations.

Advisor :: Dr. Jack Stokes, Principal RSDE @ MSR Redmond

Beijing Institute for General Artificial Intelligence (BIGAI)

Research Intern @ Natural Language and Conversational AI Lab

Jan, 2024 – May, 2024

Worked on red teaming LLMs to increase sample diversity using constrained on-policy reinforcement learning and reward shaping.

Advisor :: Dr. Zilong Zheng, Research Scientist @ BIGAI

Academic Talks

May–August, 2025: **10 Invited Talks** at Nvidia/Hugging Face/Amazon/Microsoft/ByteDance/BAAI/Gaoling School/ERIC Lab/The Network School/Plutos for **Absolute Zero Reasoner**

February, 2025: **AAAI 2025 Oral Presentation** DiveR-CT

January, 2025: **AI TIME Youth PhD Talk** - DiveR-CT

August, 2024: **IJCAI 2024 AI4Research Workshop** Invited Paper Talk for ExpeL

February, 2024: **AAAI 2024 Oral Presentation** of ExpeL: LLM Agents are Experiential Learners

January, 2024: **AI TIME Youth PhD Talk** - ExpeL: LLM Agents are Experiential Learners

February, 2023: **AI TIME Youth PhD Talk** - A Mixture of Surprises for Unsupervised Reinforcement Learning

Academic Services/Awards

2025 Oct.: Top Reviewer in NeurIPS 2025

2023–present: **Academic Reviewer** in NeurIPS 2023/2024/2025, ICLR 2024/2025, AAAI 2025/2026, ICML 2024/2025, AISTATS 2025

Last updated: January 12, 2026

Complete publication list and project pages available at
andrewzh112.github.io